

# Cambridge AI+

Lecture 1 (Part 2): Nearest Neighbour

Thomas Sauerwald

University of Cambridge, Department of Computer Science  
email: [thomas.sauerwald@cl.cam.ac.uk](mailto:thomas.sauerwald@cl.cam.ac.uk)

Nearest Neighbour Algorithm

Additional Material

## The Intuition behind Nearest Neighbour Algorithm

### Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the most similar data points and make prediction based on these classes.

# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏

# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **color** (from yellow to red)



# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **color** (from yellow to red) and **weight**




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**






# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍎
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)



- The unknown fruit is heavier and more red-ish than all seen banana
- It has the (average) colour of an orange, but it is heavier
- It has the (average) weight of an apple, but it is more yellow-ish
- Should it be an apple or orange?




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the most similar data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)




# The Intuition behind Nearest Neighbour Algorithm

## Idea of Nearest Neighbour

For any new (unseen) data point to be classified, find the **most similar** data points and make prediction based on these classes.

- **Example:** We want to classify fruits into 🍌, 🍊 and 🍏
- We will measure the **colour** (from yellow to red) and **weight**
- We now see a fruit with colour 6 (orange) and weight 6 (quite heavy)



## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	−

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	−
Training Point 3	0.58	+
Training Point 4	0.2	−



## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	—
Training Point 3	0.58	+
Training Point 4	0.2	—
Training Point 5	0.7	—

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	—
Training Point 3	0.58	+
Training Point 4	0.2	—
Training Point 5	0.7	—
Training Point 6	0.84	+
Training Point 7	0.1	+

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-

## Towards the $k$ -NN: A Simple Example in One Dimension


Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	<u>0.16</u>	??

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??





## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??



## Towards the $k$ -NN: A Simple Example in One Dimension


Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??






## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??




## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
<u>Test Point 11</u>	0.16	??



## Towards the $k$ -NN: A Simple Example in One Dimension


Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??



**Quiz 1:** How would you classify point 11?

## Towards the $k$ -NN: A Simple Example in One Dimension


Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??



**Quiz 1:** How would you classify point 11?

## Towards the $k$ -NN: A Simple Example in One Dimension


Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	??



**Quiz 1:** How would you classify point 11?

## Towards the $k$ -NN: A Simple Example in One Dimension

Input	Weight	Apple or Orange
Training Point 1	0.4	+
Training Point 2	0.35	-
Training Point 3	0.58	+
Training Point 4	0.2	-
Training Point 5	0.7	-
Training Point 6	0.84	+
Training Point 7	0.1	+
Training Point 8	0.75	+
Training Point 9	0.25	-
Training Point 10	0.05	-
Test Point 11	0.16	-




**Quiz 1:** How would you classify point 11?



## The Setup (A bit more formal...)


- Let  $\mathcal{X} = \mathbb{R}^d$  (domain set)
- Let  $\mathcal{Y} = \{-1, +1\}$  (label set)
- Let  $S = ((\underline{\mathbf{x}}_1, \underline{y}_1), (\underline{\mathbf{x}}_2, \underline{y}_2), \dots, (\underline{\mathbf{x}}_m, \underline{y}_m))$  (training set)





## The Setup (A bit more formal...)


- Let  $\mathcal{X} = \mathbb{R}^d$  (domain set)
- Let  $\mathcal{Y} = \{-1, +1\}$  (label set)
- Let  $S = ((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m))$  (training set)



**Goal:** Find a predictor  $h : \mathcal{X} \rightarrow \mathcal{Y}$ , which labels any unseen data point.

## The Setup (A bit more formal...)

- Let  $\mathcal{X} = \mathbb{R}^d$  (domain set)
- Let  $\mathcal{Y} = \{-1, +1\}$  (label set)
- Let  $S = ((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m))$  (training set)




**Goal:** Find a predictor  $h : \mathcal{X} \rightarrow \mathcal{Y}$ , which labels any unseen data point.

sometimes also called classifier or prediction rule

## The Setup (A bit more formal...)

- Let  $\mathcal{X} = \mathbb{R}^d$  (domain set)
- Let  $\mathcal{Y} = \{-1, +1\}$  (label set)
- Let  $S = ((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m))$  (training set)

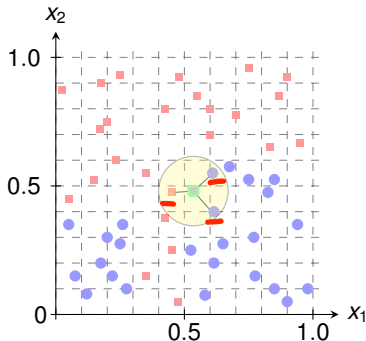


**Goal:** Find a predictor  $h : \mathcal{X} \rightarrow \mathcal{Y}$ , which labels any unseen data point.

sometimes also called classifier or prediction rule

## The Setup (A bit more formal...)

- Let  $\mathcal{X} = \mathbb{R}^d$  (domain set)
- Let  $\mathcal{Y} = \{-1, +1\}$  (label set)
- Let  $S = ((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m))$  (training set)



**Goal:** Find a predictor  $h : \mathcal{X} \rightarrow \mathcal{Y}$ , which labels any unseen data point.

sometimes also called classifier or prediction rule


## The $k$ -NN Algorithm

---

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the Euclidean Distance


# The $k$ -NN Algorithm

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the Euclidean Distance



# The $k$ -NN Algorithm

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the Euclidean Distance



## The $k$ -NN Algorithm

---

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the Euclidean Distance ✓
- For every  $\mathbf{x} \in \mathcal{X}$ , let  $\pi_1(\mathbf{x}), \dots, \pi_m(\mathbf{x})$  be a reordering of  $\{1, 2, \dots, m\}$  according to their distance to  $\mathbf{x}$ , i.e., for every  $1 \leq i < m$ :

$$\underline{d(\mathbf{x}, x_{\pi_i(\mathbf{x})})} \leq \underline{d(\mathbf{x}, x_{\pi_{i+1}(\mathbf{x})})}.$$



## The $k$ -NN Algorithm

---

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the **Euclidean Distance**
- For every  $\mathbf{x} \in \mathcal{X}$ , let  $\pi_1(\mathbf{x}), \dots, \pi_m(\mathbf{x})$  be a reordering of  $\{1, 2, \dots, m\}$  according to their distance to  $\mathbf{x}$ , i.e., for every  $1 \leq i < m$ :

$$d(\mathbf{x}, x_{\pi_i(\mathbf{x})}) \leq d(\mathbf{x}, x_{\pi_{i+1}(\mathbf{x})}).$$

For each point  $\mathbf{x}$ , we **sort** the elements in training set increasingly in their distance to  $\mathbf{x}$ .

## The $k$ -NN Algorithm

- Let  $d(\mathbf{x}, \mathbf{x}') := \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$  be the **Euclidean Distance**
- For every  $\mathbf{x} \in \mathcal{X}$ , let  $\pi_1(\mathbf{x}), \dots, \pi_m(\mathbf{x})$  be a reordering of  $\{1, 2, \dots, m\}$  according to their distance to  $\mathbf{x}$ , i.e., for every  $1 \leq i < m$ :

$$d(\mathbf{x}, x_{\pi_i(\mathbf{x})}) \leq d(\mathbf{x}, x_{\pi_{i+1}(\mathbf{x})}).$$

For each point  $\mathbf{x}$ , we **sort** the elements in training set increasingly in their distance to  $\mathbf{x}$ .

~~$k$ -NN~~

**input:** a training sample  $S = (\underline{\mathbf{x}}_1, \underline{y}_1), \dots, (\underline{\mathbf{x}}_m, \underline{y}_m)$

**output:** for every point  $\underline{\mathbf{x}} \in \mathcal{X}$ ,



return the majority label among  $\{y_{\pi_i(\mathbf{x})} : i \leq k\}$

Source: SS&BD

role  $k$ !

## Special Case: 1-NN

$$K=1, d=2$$




Source: Zemel, Urtasun, Fidler

- For  $k = 1$ , the produced decision boundaries are Voronoi-cells
- Any new point will be classified according to the centre of each cell

## How many Neighbours should we choose?



$k = 1$




$k = 15$


Source: Lecture by Ulrike von Luxburg

## How many Neighbours should we choose?

For small  $k$ ,  $k$ -NN overfits the data!



$k = 1$




$k = 15$

Source: Lecture by Ulrike von Luxburg


## How many Neighbours should we choose?

For small  $k$ ,  $k$ -NN overfits the data!



$k = 1$

Rule-of-thumb:  $k \approx \log m$  is good




$k = 15$

Source: Lecture by Ulrike von Luxburg


## How many Neighbours should we choose?

For small  $k$ ,  $k$ -NN overfits the data!

Rule-of-thumb:  $k \approx \log m$  is good



$k = 1$




$k = 15$

Source: Lecture by Ulrike von Luxburg

Quiz 2: What happens if  ~~$k = m$~~ , where  $m$  is the total number of points in our training set?




## Training Error, Test Error and Overfitting







# Training Error, Test Error and Overfitting



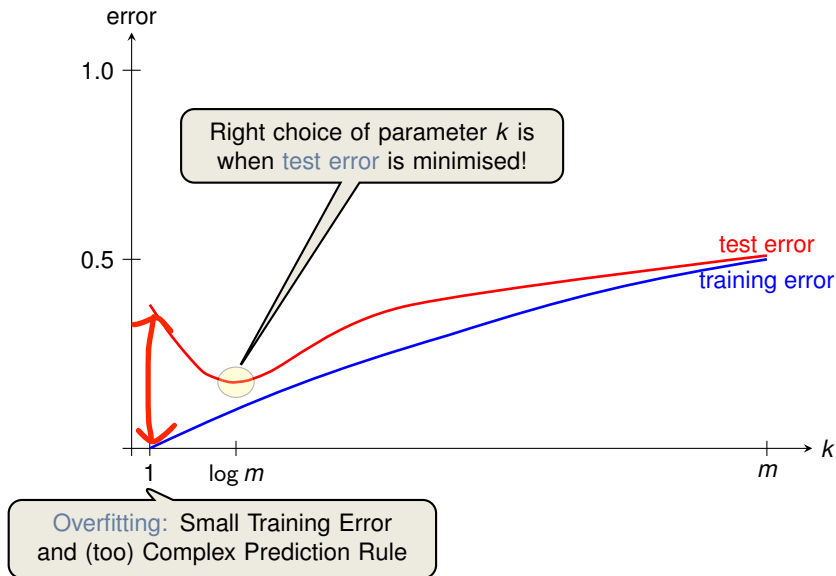
# Training Error, Test Error and Overfitting



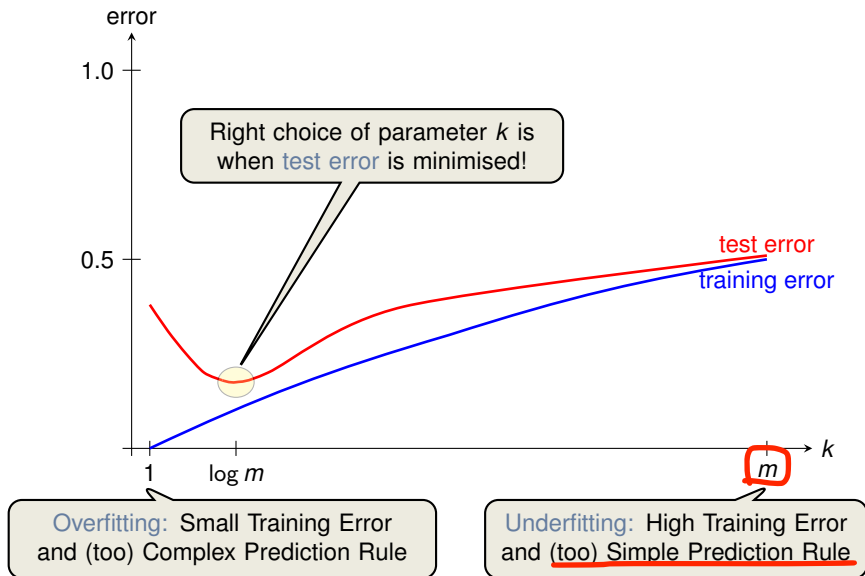
## Training Error, Test Error and Overfitting



## Training Error, Test Error and Overfitting



## Training Error, Test Error and Overfitting



## Feature Weighting and Distance Weighting

---

### Feature Weighting

- k-NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ Normalise all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  **nearest neighbour** without taking **distance** into account



## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$
  2. **Inverse Distance**:  $i$ -th closest point is weighted by  $1 / \text{dist}(\mathbf{x}_i, \mathbf{x})$

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$
  2. **Inverse Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})$
  3. **Inverse Square Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})^2$

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$
  2. **Inverse Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})$
  3. **Inverse Square Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})^2$
- Alternatively, one can take **Gaussian Weighting** over all points: each point  $\mathbf{x}_i$  is weighted by  $\exp(-c \cdot (\mathbf{x}_i - \mathbf{x})^2)$

## Feature Weighting and Distance Weighting

### Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**

### Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$
  2. **Inverse Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})$
  3. **Inverse Square Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})^2$
- Alternatively, one can take **Gaussian Weighting** over all points: each point  $\mathbf{x}_i$  is weighted by  $\exp(-c \cdot (\mathbf{x}_i - \mathbf{x})^2)$

No need to choose  $k$ , but  $c$  controls sensitivity of nearby points.

# Feature Weighting and Distance Weighting

## Feature Weighting

- $k$ -NN is not **scaling invariant**, i.e., multiplying the the same feature of each data point may affect the results
- ⇒ **Normalise** all features so that the **square deviation** from **mean** is the same for all  $d$  **features**



## Distance Weighting

- In  $k$ -NN, we simply take **majority** over the  $k$  nearest neighbour without taking **distance** into account
- There are several **weighting-schemes**:
  1. **Harmonic Weighting**:  $i$ -th closest point is weighted by  $1/i$
  2. **Inverse Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})$
  3. **Inverse Square Distance**:  $i$ -th closest point is weighted by  $1/\text{dist}(\mathbf{x}_i, \mathbf{x})^2$
- Alternatively, one can take **Gaussian Weighting** over all points: each point  $\mathbf{x}_i$  is weighted by  $\exp(-c \cdot (\mathbf{x}_i - \mathbf{x})^2)$

No need to choose  $k$ , but  $c$  controls sensitivity of nearby points.

We see some examples of **Gaussian Weighting** in next lecture!

# Uniform Weighting vs. Inverse Distance Weighting





Source: <https://scikit-learn.org/stable/modules/neighbors.html>

smooth



# Uniform Weighting vs. Inverse Distance Weighting



Source: <https://scikit-learn.org/stable/modules/neighbors.html>

Distance-Weighting usually produces smoother decision boundaries and is less likely to overfit (or underfit).


## Summary of Nearest Neighbour

---

- + very easy to implement and understand
- + does not use any training/learning phase
- + can be applied to almost any prediction problem  
(often a good “baseline”)

## Summary of Nearest Neighbour

---

- + very easy to implement and understand
- + does not use any training/learning phase
- + can be applied to almost any prediction problem (often a good “baseline”)
- results may crucially depend on the choice of distance function, representation of features and  $k$   Log m
- finding nearest neighbour is costly

## Summary of Nearest Neighbour




---

- + very easy to implement and understand
- + does not use any training/learning phase
- + can be applied to almost any prediction problem (often a good “baseline”)
- results may crucially depend on the choice of distance function, representation of features and  $k$
- finding nearest neighbour is costly






in practice, we often use a powerful pre-processing tool called **Dimensionality Reduction!**

## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)






## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)



- For the computer, each image is just a long sequence of bits


...

## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)



- For the computer, each image is just a long sequence of bits  
Example: the first image may start 0, 0, 0, 0, 1, 1, 1, 0, 0, ...


## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)



- For the computer, each image is just a long sequence of bits  
Example: the first image may start 0, 0, 0, 0, 1, 1, 1, 0, 0, ...
- For a  $1024 \times 768$  image, we have 786432 pixels




## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)



- For the computer, each image is just a long sequence of bits  
Example: the first image may start 0, 0, 0, 0, 1, 1, 1, 0, 0, ...
- For a  $1024 \times 768$  image, we have 786432 pixels
- if coloured, we have 3 RGB values in the range  $[0, 255]$  (8 bits)

## Exercise 2: Applying $k$ -NN to Character Recognition (1/2)



- For the computer, each image is just a long sequence of bits  
Example: the first image may start 0, 0, 0, 0, 1, 1, 1, 0, 0, ...
- For a  $1024 \times 768$  image, we have 786432 pixels
- if coloured, we have 3 RGB values in the range  $[0, 255]$  (8 bits)
- Overall the picture is represented  $786432 \cdot 3 \cdot 8 = 18,874,368$  bits

## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

---


## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

---


Apply 2-NN with distance being the number of differently coloured cells.

## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)


Apply 2-NN with distance being the number of differently coloured cells.




—



—




+




+

## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)


Apply 2-NN with distance being the number of differently coloured cells.




—




—



+




+



??


## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.



## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)


Apply 2-NN with distance being the number of differently coloured cells.






## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.




## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.




## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.




## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.




## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.




## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.



## Exercise 2: Applying $k$ -NN to Character Recognition (2/2)

Apply 2-NN with distance being the number of differently coloured cells.





A. Blum, J. Hopcroft and R. Kannan  
Foundations of Data Science  
1st edition, Cambridge University Press, 2020.  
<https://www.cs.cornell.edu/jeh/book.pdf>



J. Leskovec, A. Rajaraman, J. D. Ullmann  
Mining of Massive Datasets.  
3rd edition, Cambridge University Press, 2020.



S. Shalev-Shwartz and S. Ben-David  
Understanding Machine Learning: From Theory to Algorithms  
Cambridge University Press, 2014.  
[https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/  
understanding-machine-learning-theory-algorithms.pdf](https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/understanding-machine-learning-theory-algorithms.pdf)



Nearest Neighbour Algorithm

Additional Material

### Algorithm

- Input:
  - Rating matrix with  $n$  users and  $p$  movies,
  - integer  $k \geq 1$ ,
  - unknown rating  $x_{i,j}$  (to be predicted by algorithm)
- Find a set of  $k$  movies  $M$  most similar to item  $j$  that are rated by user  $i$
- Output:

$$x_{i,j} = \frac{\sum_{\ell \in M} \text{sim}(j, \ell) \cdot x_{i,\ell}}{\sum_{\ell \in M} \text{sim}(j, \ell)}$$

## Some Ideas on efficient Implementation

---

- $m$  number of points,  $d$  dimension, **want:**  $k$  nearest neighbours to input  $x$

## Some Ideas on efficient Implementation

---

- $m$  number of points,  $d$  dimension, **want:**  $k$  nearest neighbours to input  $x$

### Basic Algorithm

- Compute all distances from the  $m$  points to  $x$
- Scan list for the nearest neighbour of  $x$  and then remove it from list
- Running Time is  $O(m \cdot d + m \cdot k)$

## Some Ideas on efficient Implementation

- $m$  number of points,  $d$  dimension, **want:**  $k$  nearest neighbours to input  $x$

### Basic Algorithm

- Compute all distances from the  $m$  points to  $x$
- Scan list for the nearest neighbour of  $x$  and then remove it from list
- Running Time is  $O(m \cdot d + m \cdot k)$

### Algorithm based on Sorting

- Compute all distances from the  $m$  points to  $x$
- **Sort** all these  $m$  distances increasingly
- Running Time is  $O(m \cdot d + m \cdot \log m)$

## Some Ideas on efficient Implementation

- $m$  number of points,  $d$  dimension, **want**:  $k$  nearest neighbours to input  $x$

### Basic Algorithm

- Compute all distances from the  $m$  points to  $x$
- Scan list for the nearest neighbour of  $x$  and then remove it from list
- Running Time is  $O(m \cdot d + m \cdot k)$


### Algorithm based on Sorting

- Compute all distances from the  $m$  points to  $x$
- **Sort** all these  $m$  distances increasingly
- Running Time is  $O(m \cdot d + m \cdot \log m)$

### Clever Algorithm


- Compute all distances from the  $m$  points to  $x$
- **Use Quick-Select** to find the  $k$  nearest distances (without sorting)
- Running Time is  $O(m \cdot d + m)$  if we want the list of  $k$  nearest points, and  $O(m \cdot d + m + k \log k)$  if we want the  $k$  nearest points in order.

# Curse of Dimensionality



Source: Kilian Weinberger

# Curse of Dimensionality




Source: Kilian Weinberger

- Suppose  $m = 1000$  points are “randomly” spread across  $[0, 1]^d$




# Curse of Dimensionality



Source: Kilian Weinberger

- Suppose  $m = 1000$  points are “randomly” spread across  $[0, 1]^d$
- If we want to find the 10 nearest neighbour, how large must be the subcube be?

# Curse of Dimensionality




Source: Kilian Weinberger

- Suppose  $m = 1000$  points are “randomly spread” across  $[0, 1]^d$
- If we want to find the 10 nearest neighbour, how large must be the subcube be?

$d$	$\ell$
2	0.1
10	0.63
100	0.955
1000	0.9954

# Curse of Dimensionality




Source: Kilian Weinberger

- Suppose  $m = 1000$  points are “randomly spread” across  $[0, 1]^d$
- If we want to find the 10 nearest neighbour, how large must be the subcube be?

$d$	$\ell$
2	0.1
10	0.63
100	0.955
1000	0.9954

To find the closest neighbour, we need to search the entire space!

# Curse of Dimensionality



Source: Kilian Weinberger

In high dimensions, almost all points have the same (far) distance!

- Suppose  $m = 1000$  points are “randomly” spread across  $[0, 1]^d$
- If we want to find the 10 nearest neighbour, how large must be the subcube be?

$d$	$\ell$
2	0.1
10	0.63
100	0.955
1000	0.9954

To find the closest neighbour, we need to search the entire space!